

# FacialRecognition

January 4, 2021

## 1 Facial Recognition using CNN

```
[1]: import numpy as np
import pandas as pd
from skimage import io
from keras.utils import to_categorical
import matplotlib.pyplot as plt
%matplotlib inline
```

Using TensorFlow backend.

### 1.1 Data Pre-processing


```
[2]: dt=np.load('ORL_faces.npz',allow_pickle=True)
trainX=dt["trainX"].reshape(240,112,92)
trainY=dt["trainY"]
testX=dt["testX"].reshape(160,112,92)
testY=dt["testY"]
```

```
[3]: trainX=trainX.reshape(-1,112,92,1)
testX=testX.reshape(-1,112,92,1)

#Normalizing the data
trainX=trainX.astype('float32')/225
testX=testX.astype('float32')/255

#encoding the targets to categorical value
ohe_trainY=to_categorical(trainY)
ohe_testY=to_categorical(testY)
```

```
[100]: for i in range(1,12):
        plt.subplot(4,3,i)
        plt.imshow(trainX[i])
```



```
[80]: print(trainX.shape, ohe_trainY.shape, testX.shape, ohe_testY.shape)
```

```
(240, 112, 92, 1) (240, 20) (160, 112, 92, 1) (160, 20)
```

## 1.2 Model Architecting

```
[4]: import keras
from keras.losses import categorical_crossentropy
from keras.optimizers import Adam
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten
from keras.layers import Conv2D, MaxPooling2D
from keras.layers.normalization import BatchNormalization
from keras.layers.advanced_activations import LeakyReLU
```

```
[5]: model=Sequential()
model.
    ↳ add(Conv2D(32, kernel_size=(3, 3), activation='linear', input_shape=(112, 92, 1), padding='same'))
model.add(LeakyReLU(alpha=0.05))
model.add(MaxPooling2D((2, 2), padding='same'))
model.add(Dropout(0.25))

model.add(Flatten())
model.add(Dense(128, activation='linear'))
```

```

model.add(LeakyReLU(alpha=0.05))
model.add(Dropout(0.5))

model.add(Dense(20,activation='softmax'))
model.
    ↪ compile(loss='categorical_crossentropy',optimizer=Adam(),metrics=['accuracy'])
model.summary()
model_train=model.
    ↪ fit(trainX,ohe_trainY,batch_size=128,epochs=50,verbose=1,validation_data=(testX,ohe_testY))

```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 112, 92, 32)	320
leaky_re_lu_1 (LeakyReLU)	(None, 112, 92, 32)	0
max_pooling2d_1 (MaxPooling2D)	(None, 56, 46, 32)	0
dropout_1 (Dropout)	(None, 56, 46, 32)	0
flatten_1 (Flatten)	(None, 82432)	0
dense_1 (Dense)	(None, 128)	10551424
leaky_re_lu_2 (LeakyReLU)	(None, 128)	0
dropout_2 (Dropout)	(None, 128)	0
dense_2 (Dense)	(None, 20)	2580

Total params: 10,554,324

Trainable params: 10,554,324

Non-trainable params: 0

Train on 240 samples, validate on 160 samples

Epoch 1/50

240/240 [=====] - 2s 10ms/step - loss: 7.6039 - accuracy: 0.0625 - val\_loss: 5.4096 - val\_accuracy: 0.0500

Epoch 2/50

240/240 [=====] - 2s 10ms/step - loss: 8.1429 - accuracy: 0.0542 - val\_loss: 3.4554 - val\_accuracy: 0.0500

Epoch 3/50

240/240 [=====] - 2s 10ms/step - loss: 4.5085 - accuracy: 0.0833 - val\_loss: 2.9383 - val\_accuracy: 0.1125

Epoch 4/50

240/240 [=====] - 2s 10ms/step - loss: 3.2116 - accuracy: 0.1417 - val\_loss: 2.9717 - val\_accuracy: 0.1125  
Epoch 5/50  
240/240 [=====] - 2s 10ms/step - loss: 3.1104 - accuracy: 0.1042 - val\_loss: 2.8462 - val\_accuracy: 0.1937  
Epoch 6/50  
240/240 [=====] - 2s 10ms/step - loss: 3.0706 - accuracy: 0.0958 - val\_loss: 2.6564 - val\_accuracy: 0.2562  
Epoch 7/50  
240/240 [=====] - 2s 10ms/step - loss: 2.8574 - accuracy: 0.1750 - val\_loss: 2.5908 - val\_accuracy: 0.2438  
Epoch 8/50  
240/240 [=====] - 2s 10ms/step - loss: 2.6832 - accuracy: 0.1708 - val\_loss: 2.4722 - val\_accuracy: 0.3438  
Epoch 9/50  
240/240 [=====] - 2s 10ms/step - loss: 2.5729 - accuracy: 0.1750 - val\_loss: 2.3625 - val\_accuracy: 0.2937  
Epoch 10/50  
240/240 [=====] - 2s 10ms/step - loss: 2.6275 - accuracy: 0.1583 - val\_loss: 2.3338 - val\_accuracy: 0.3500  
Epoch 11/50  
240/240 [=====] - 2s 10ms/step - loss: 2.3134 - accuracy: 0.2625 - val\_loss: 2.2884 - val\_accuracy: 0.2375  
Epoch 12/50  
240/240 [=====] - 2s 10ms/step - loss: 2.3009 - accuracy: 0.2583 - val\_loss: 2.1267 - val\_accuracy: 0.5437  
Epoch 13/50  
240/240 [=====] - 2s 10ms/step - loss: 2.1500 - accuracy: 0.3167 - val\_loss: 2.0365 - val\_accuracy: 0.5125  
Epoch 14/50  
240/240 [=====] - 2s 10ms/step - loss: 2.0901 - accuracy: 0.3292 - val\_loss: 1.8760 - val\_accuracy: 0.7063  
Epoch 15/50  
240/240 [=====] - 2s 10ms/step - loss: 1.8609 - accuracy: 0.4083 - val\_loss: 1.8308 - val\_accuracy: 0.7188  
Epoch 16/50  
240/240 [=====] - 2s 10ms/step - loss: 1.7804 - accuracy: 0.4292 - val\_loss: 1.6782 - val\_accuracy: 0.7063  
Epoch 17/50  
240/240 [=====] - 2s 10ms/step - loss: 1.6483 - accuracy: 0.4750 - val\_loss: 1.5921 - val\_accuracy: 0.6313  
Epoch 18/50  
240/240 [=====] - 2s 10ms/step - loss: 1.5644 - accuracy: 0.4958 - val\_loss: 1.3564 - val\_accuracy: 0.6125  
Epoch 19/50  
240/240 [=====] - 2s 10ms/step - loss: 1.2800 - accuracy: 0.6292 - val\_loss: 1.2799 - val\_accuracy: 0.8750  
Epoch 20/50

240/240 [=====] - 2s 10ms/step - loss: 1.2190 - accuracy: 0.6208 - val\_loss: 1.1918 - val\_accuracy: 0.9062  
Epoch 21/50  
240/240 [=====] - 2s 10ms/step - loss: 1.1799 - accuracy: 0.6625 - val\_loss: 1.0644 - val\_accuracy: 0.9187  
Epoch 22/50  
240/240 [=====] - 2s 10ms/step - loss: 0.9595 - accuracy: 0.7417 - val\_loss: 0.9392 - val\_accuracy: 0.9250  
Epoch 23/50  
240/240 [=====] - 2s 10ms/step - loss: 1.0413 - accuracy: 0.6708 - val\_loss: 0.8458 - val\_accuracy: 0.9125  
Epoch 24/50  
240/240 [=====] - 2s 10ms/step - loss: 0.8513 - accuracy: 0.7542 - val\_loss: 0.8066 - val\_accuracy: 0.9187  
Epoch 25/50  
240/240 [=====] - 2s 10ms/step - loss: 0.7066 - accuracy: 0.7792 - val\_loss: 0.7795 - val\_accuracy: 0.9000  
Epoch 26/50  
240/240 [=====] - 2s 10ms/step - loss: 0.6831 - accuracy: 0.7833 - val\_loss: 0.6447 - val\_accuracy: 0.9125  
Epoch 27/50  
240/240 [=====] - 2s 10ms/step - loss: 0.5509 - accuracy: 0.8292 - val\_loss: 0.5895 - val\_accuracy: 0.9250  
Epoch 28/50  
240/240 [=====] - 2s 10ms/step - loss: 0.5535 - accuracy: 0.8458 - val\_loss: 0.5096 - val\_accuracy: 0.9375  
Epoch 29/50  
240/240 [=====] - 2s 10ms/step - loss: 0.4282 - accuracy: 0.8625 - val\_loss: 0.5721 - val\_accuracy: 0.9250  
Epoch 30/50  
240/240 [=====] - 2s 10ms/step - loss: 0.4993 - accuracy: 0.8750 - val\_loss: 0.5776 - val\_accuracy: 0.9500  
Epoch 31/50  
240/240 [=====] - 2s 10ms/step - loss: 0.3430 - accuracy: 0.9250 - val\_loss: 0.5679 - val\_accuracy: 0.9312  
Epoch 32/50  
240/240 [=====] - 2s 10ms/step - loss: 0.4079 - accuracy: 0.8917 - val\_loss: 0.5503 - val\_accuracy: 0.9187  
Epoch 33/50  
240/240 [=====] - 2s 10ms/step - loss: 0.3459 - accuracy: 0.8875 - val\_loss: 0.4246 - val\_accuracy: 0.9125  
Epoch 34/50  
240/240 [=====] - 2s 10ms/step - loss: 0.2627 - accuracy: 0.9167 - val\_loss: 0.4148 - val\_accuracy: 0.9250  
Epoch 35/50  
240/240 [=====] - 2s 10ms/step - loss: 0.2964 - accuracy: 0.9333 - val\_loss: 0.4261 - val\_accuracy: 0.9062  
Epoch 36/50

```

240/240 [=====] - 2s 10ms/step - loss: 0.2420 -
accuracy: 0.9333 - val_loss: 0.3981 - val_accuracy: 0.9187
Epoch 37/50
240/240 [=====] - 2s 10ms/step - loss: 0.2546 -
accuracy: 0.9250 - val_loss: 0.3636 - val_accuracy: 0.9625
Epoch 38/50
240/240 [=====] - 2s 10ms/step - loss: 0.1555 -
accuracy: 0.9792 - val_loss: 0.3488 - val_accuracy: 0.9625
Epoch 39/50
240/240 [=====] - 2s 10ms/step - loss: 0.1837 -
accuracy: 0.9667 - val_loss: 0.3377 - val_accuracy: 0.9438
Epoch 40/50
240/240 [=====] - 2s 10ms/step - loss: 0.1760 -
accuracy: 0.9625 - val_loss: 0.3344 - val_accuracy: 0.9500
Epoch 41/50
240/240 [=====] - 2s 10ms/step - loss: 0.1369 -
accuracy: 0.9625 - val_loss: 0.3438 - val_accuracy: 0.9375
Epoch 42/50
240/240 [=====] - 2s 10ms/step - loss: 0.1678 -
accuracy: 0.9708 - val_loss: 0.3335 - val_accuracy: 0.9500
Epoch 43/50
240/240 [=====] - 2s 10ms/step - loss: 0.1388 -
accuracy: 0.9667 - val_loss: 0.3441 - val_accuracy: 0.9375
Epoch 44/50
240/240 [=====] - 2s 10ms/step - loss: 0.1328 -
accuracy: 0.9875 - val_loss: 0.3293 - val_accuracy: 0.9625
Epoch 45/50
240/240 [=====] - 2s 10ms/step - loss: 0.1030 -
accuracy: 0.9792 - val_loss: 0.3175 - val_accuracy: 0.9563
Epoch 46/50
240/240 [=====] - 2s 10ms/step - loss: 0.1355 -
accuracy: 0.9583 - val_loss: 0.3218 - val_accuracy: 0.9375
Epoch 47/50
240/240 [=====] - 2s 10ms/step - loss: 0.1046 -
accuracy: 0.9875 - val_loss: 0.3243 - val_accuracy: 0.9312
Epoch 48/50
240/240 [=====] - 2s 10ms/step - loss: 0.1051 -
accuracy: 0.9708 - val_loss: 0.3309 - val_accuracy: 0.9500
Epoch 49/50
240/240 [=====] - 2s 10ms/step - loss: 0.0886 -
accuracy: 0.9917 - val_loss: 0.3356 - val_accuracy: 0.9312
Epoch 50/50
240/240 [=====] - 2s 10ms/step - loss: 0.1211 -
accuracy: 0.9667 - val_loss: 0.3089 - val_accuracy: 0.9500

```

```

[8]: # Model Evaluation
test_eval=model.evaluate(testX,ohe_testY,verbose=1)

```


```


print("test loss : ",test_eval[0])
print("test accuracy : ",test_eval[1])
accuracy=model_train.history['accuracy']
val_accuracy=model_train.history['val_accuracy']
loss=model_train.history['loss']
val_loss=model_train.history['val_loss']
epochs=range(len(accuracy))
plt.plot(epochs,accuracy,'bo',label='training accuracy')
plt.plot(epochs,val_accuracy,'b',label='validation accuracy')
plt.title('training and validation accuracy')
plt.legend()
plt.figure()
plt.plot(epochs,loss,'bo',label='training accuracy')
plt.plot(epochs,val_loss,'b',label='validation accuracy')
plt.title('training and validation loss')
plt.legend()
plt.figure()

```

160/160 [=====] - 0s 2ms/step  
test loss : 0.3089377969503403  
test accuracy : 0.949999988079071

[8]: <Figure size 432x288 with 0 Axes>





<Figure size 432x288 with 0 Axes>